

ML computer_image_classification

May 28, 2019

0.0.1 Computer vision for image classification

Problem Statement

After cleaning the dataset, work on the classification

The instances were drawn randomly from a database of 7 outdoor images. The images were hand-segmented to create a classification for every pixel. Each instance is a 3 pixel x 3 pixel region (9 pixels).

Your task is to predict the probability that each 3x3 image region belongs to each of the seven classes (window, foliage, brickface, sky, grass, path, cement).

For each unique 3x3 region you should provide 7 probabilities with value between 0 and 1.

0.0.2 import packages

```
In [1]: %matplotlib inline
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import scipy as sc
import sklearn as skl
import seaborn as sns
```

```
In [2]: # Input data files are available in the "../input/" directory.
# For example, running this (by clicking run or pressing Shift+Enter) will list the fi
```

```
import os
print(os.listdir("../cmu_data"))
```

```
['.ipynb_checkpoints', 'train.csv', 'clean_image.csv', 'sample_submission_fixed.csv', 'test.csv']
```

0.0.3 import dataset

```
In [3]: #import csv data and read it as a dataframe
image_test_df=pd.read_csv('../cmu_data/test.csv', sep=',')
image_train_df=pd.read_csv('../cmu_data/train.csv', sep=',')
```

0.1 EDA

```
In [4]: image_test_df.head(2)
```

```
Out [4]:
```

	id	REGION-CENTROID-COL	REGION-CENTROID-ROW	REGION-PIXEL-COUNT	\
0	0	105	110	9	
1	1	36	189	9	

	SHORT-LINE-DENSITY-5	SHORT-LINE-DENSITY-2	VEDGE-MEAN	VEDGE-SD	\
0	0.0	0.0	0.500000	0.122222	
1	0.0	0.0	1.944444	2.462961	

	HEDGE-MEAN	HEDGE-SD	INTENSITY-MEAN	RAWRED-MEAN	RAWBLUE-MEAN	\
0	1.944444	2.374073	4.222222	4.333334	6.333334	
1	5.833334	21.588884	31.370370	28.444445	38.000000	

	RAWGREEN-MEAN	EXRED-MEAN	EXBLUE-MEAN	EXGREEN-MEAN	VALUE-MEAN	\
0	2.000000	0.333333	6.333334	-6.666666	6.333334	
1	27.666666	-8.777778	19.888890	-11.111111	38.000000	

	SATURATION-MEAN	HUE-MEAN
0	0.708333	-1.539762
1	0.266302	-2.020345

```
In [5]: image_train_df.head(2)
```

```
Out [5]:
```

	REGION-CENTROID-COL	REGION-CENTROID-ROW	REGION-PIXEL-COUNT	\
0	123.0	152.0	9	
1	226.0	110.0	9	

	SHORT-LINE-DENSITY-5	SHORT-LINE-DENSITY-2	VEDGE-MEAN	VEDGE-SD	\
0	0.0	0.0	0.000000	0.000000	
1	0.0	0.0	0.333333	0.088889	

	HEDGE-MEAN	HEDGE-SD	INTENSITY-MEAN	RAWRED-MEAN	RAWBLUE-MEAN	\
0	0.0	0.000000	0.000000	0.000000	0.000000	
1	0.5	0.211111	1.666667	0.111111	4.444445	

	RAWGREEN-MEAN	EXRED-MEAN	EXBLUE-MEAN	EXGREEN-MEAN	VALUE-MEAN	\
0	0.000000	0.000000	0.000000	0.000000	0.000000	
1	0.444444	-4.666666	8.333333	-3.666667	4.444445	

	SATURATION-MEAN	HUE-MEAN	CLASS
0	0.000000	0.000000	0
1	0.977778	-2.155984	1

```
In [6]: image_train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 210 entries, 0 to 209
```

```
Data columns (total 20 columns):
REGION-CENTROID-COL      210 non-null float64
REGION-CENTROID-ROW      210 non-null float64
REGION-PIXEL-COUNT       210 non-null int64
SHORT-LINE-DENSITY-5     210 non-null float64
SHORT-LINE-DENSITY-2     210 non-null float64
VEDGE-MEAN                210 non-null float64
VEDGE-SD                 210 non-null float64
HEDGE-MEAN                210 non-null float64
HEDGE-SD                 210 non-null float64
INTENSITY-MEAN            210 non-null float64
RAWRED-MEAN               210 non-null float64
RAWBLUE-MEAN              210 non-null float64
RAWGREEN-MEAN             210 non-null float64
EXRED-MEAN                210 non-null float64
EXBLUE-MEAN               210 non-null float64
EXGREEN-MEAN              210 non-null float64
VALUE-MEAN                210 non-null float64
SATURATION-MEAN           210 non-null float64
HUE-MEAN                  210 non-null float64
CLASS                     210 non-null int64
dtypes: float64(18), int64(2)
memory usage: 32.9 KB
```

```
In [7]: image_train_df.shape
```

```
Out[7]: (210, 20)
```

```
In [8]: # to change dtype use .astype()
        image_train_df['REGION-PIXEL-COUNT'] = image_train_df['REGION-PIXEL-COUNT'].astype(float)
```

0.1.1 Clean datasets

```
In [9]: image_train_df.isnull().values.any()
```

```
Out[9]: False
```

```
In [10]: image_test_df[image_test_df.isnull().values.any(axis=1)]
```

```
Out[10]: Empty DataFrame
         Columns: [id, REGION-CENTROID-COL, REGION-CENTROID-ROW, REGION-PIXEL-COUNT, SHORT-LINE-DENSITY-5, SHORT-LINE-DENSITY-2, VEDGE-MEAN, VEDGE-SD, HEDGE-MEAN, HEDGE-SD, INTENSITY-MEAN, RAWRED-MEAN, RAWBLUE-MEAN, RAWGREEN-MEAN, EXRED-MEAN, EXBLUE-MEAN, EXGREEN-MEAN, VALUE-MEAN, SATURATION-MEAN, HUE-MEAN, CLASS]
         Index: []
```

0.1.2 check for outliers on train dataframe

0.1.3 IQR method

```
In [11]: import numpy as np
```

```
def outliers_iqr(image_train_df):
    quartile_1, quartile_3 = np.percentile(image_class_df, [25, 75])
    iqr = quartile_3 - quartile_1
    lower_bound = quartile_1 - (iqr * 1.5)
    upper_bound = quartile_3 + (iqr * 1.5)
    return np.where((image_train_df > upper_bound) | (image_train_df < lower_bound))
```

0.1.4 z-score method

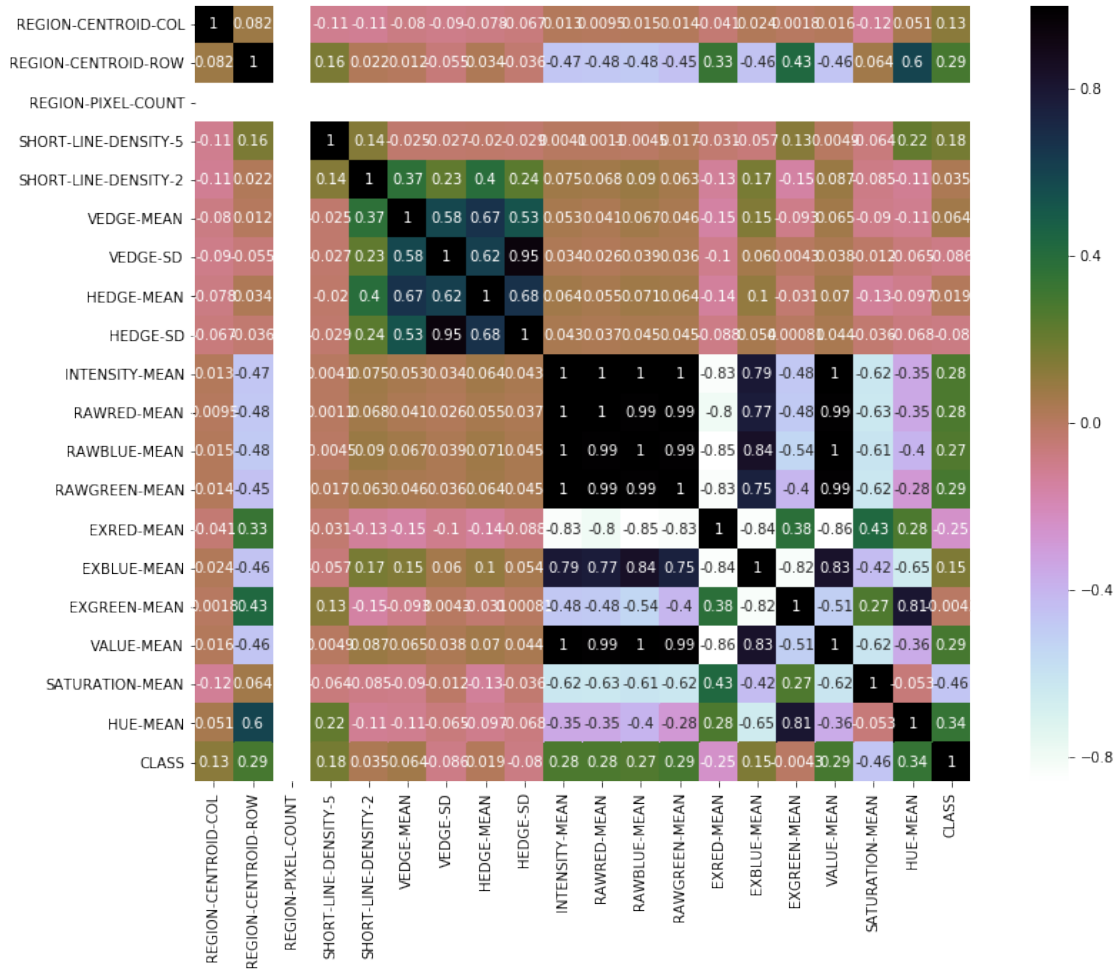
In [12]: import numpy as np

```
def outliers_z_score(image_train_df):
    threshold = 3

    mean_y = np.mean(image_train_df)
    stdev_y = np.std(image_train_df)
    z_scores = [(y - mean_y) / stdev_y for y in image_train_df]
    return np.where(np.abs(z_scores) > threshold)
```


In [13]: plt.figure(figsize=(15, 10))
sns.heatmap(image_train_df.corr(), vmax=1, annot=True, square=True, cmap='cubehelix_r')

Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x7f07c8c0dba8>



0.2 Boxplot

```
In [14]: sns.boxplot(data=image_train_df)
plt.rcParams['figure.figsize'] = (30, 20)
```



0.2.1 Plot Data

0.3 Train and test

0.3.1 train and test split

```
In [15]: image_train_df.keys()
```

```
Out[15]: Index(['REGION-CENTROID-COL', 'REGION-CENTROID-ROW', 'REGION-PIXEL-COUNT',
               'SHORT-LINE-DENSITY-5', 'SHORT-LINE-DENSITY-2', 'VEDGE-MEAN',
               'VEDGE-SD', 'HEDGE-MEAN', 'HEDGE-SD', 'INTENSITY-MEAN', 'RAWRED-MEAN',
               'RAWBLUE-MEAN', 'RAWGREEN-MEAN', 'EXRED-MEAN', 'EXBLUE-MEAN',
               'EXGREEN-MEAN', 'VALUE-MEAN', 'SATURATION-MEAN', 'HUE-MEAN', 'CLASS'],
              dtype='object')
```

```
In [16]: image_test_df.keys()
```

```
Out[16]: Index(['id', 'REGION-CENTROID-COL', 'REGION-CENTROID-ROW',
               'REGION-PIXEL-COUNT', 'SHORT-LINE-DENSITY-5', 'SHORT-LINE-DENSITY-2',
               'VEDGE-MEAN', 'VEDGE-SD', 'HEDGE-MEAN', 'HEDGE-SD', 'INTENSITY-MEAN',
               'RAWRED-MEAN', 'RAWBLUE-MEAN', 'RAWGREEN-MEAN', 'EXRED-MEAN',
               'EXBLUE-MEAN', 'EXGREEN-MEAN', 'VALUE-MEAN', 'SATURATION-MEAN',
               'HUE-MEAN'],
              dtype='object')
```

```
In [17]: # the classes
         image_train_df['CLASS'].unique()
```

```
Out[17]: array([0, 1, 2, 3, 4, 5, 6])
```

0.4 Train-Test Data Splitting

```
In [ ]: import pandas as pd
        from sklearn import datasets, linear_model
        from sklearn.model_selection import train_test_split
        from matplotlib import pyplot as plt

        columns = ['REGION-CENTROID-COL', 'REGION-CENTROID-ROW', 'REGION-PIXEL-COUNT',
                   'SHORT-LINE-DENSITY-5', 'SHORT-LINE-DENSITY-2', 'VEDGE-MEAN',
                   'VEDGE-SD', 'HEDGE-MEAN', 'HEDGE-SD', 'INTENSITY-MEAN', 'RAWRED-MEAN',
                   'RAWBLUE-MEAN', 'RAWGREEN-MEAN', 'EXRED-MEAN', 'EXBLUE-MEAN',
                   'EXGREEN-MEAN', 'VALUE-MEAN', 'SATURATION-MEAN', 'HUE-MEAN']

        holdout = image_test_df # from now on we will refer to this
                                # dataframe as the holdout data
        X = image_train_df[columns]
        y = image_train_df['CLASS']

        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20, random_state=)
```

0.4.1 Prediction of the instances with a given 3*3 pixel region

0.5 Random Forests

```
In [ ]: from sklearn.ensemble import RandomForestClassifier
        rf = RandomForestClassifier(random_state = 42)
        from pprint import pprint
        #Look at parameters used by our current forest
        print('Parameters currently in use:\n')
        pprint(rf.get_params())

In [ ]: from sklearn.model_selection import RandomizedSearchCV
        # Number of trees in random forest
        n_estimators = [int(x) for x in np.linspace(start = 100, stop = 1000, num = 10)]
        # Number of features to consider at every split
        max_features = ['auto', 'sqrt']
        # Maximum number of levels in tree
        max_depth = [int(x) for x in np.linspace(10, 110, num = 11)]
        max_depth.append(None)
        # Minimum number of samples required to split a node
        min_samples_split = [2, 5, 10]
        # Minimum number of samples required at each leaf node
        min_samples_leaf = [1, 2, 4]
        # Method of selecting samples for training each tree
        bootstrap = [True, False]
        # Create the random grid
        random_grid = {'n_estimators': n_estimators,
                       'max_features': max_features,
                       'max_depth': max_depth,
```

```

        'min_samples_split': min_samples_split,
        'min_samples_leaf': min_samples_leaf,
        'bootstrap': bootstrap}
pprint(random_grid)

In [ ]: from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier(n_estimators=800, max_depth=10, min_samples_split=2,
                           min_samples_leaf=4, max_features='sqrt', bootstrap=True, random_state=42)
rfc.fit(X_train, y_train)

In [ ]: from sklearn.preprocessing import LabelEncoder

labels = LabelEncoder()
y_train_labels_fit = labels.fit(y_train)
y_train_labels_trf = labels.transform(y_train)

x_test = holdout.drop(['id'], axis=1)
holdout_predictions = rfc.predict_proba(x_test)

# model accuracy for X_test
accuracy = rfc.score(X_test, y_test)
print (accuracy)

In [ ]: test_pred = pd.DataFrame(rfc.predict_proba(x_test)*1, columns=labels.classes_)
q = {'id': test_data["id"], 'no_financial_services': test_pred[0], 'other_only': test_pred[1],
     'mm_only': test_pred[2], 'mm_plus': test_pred[3]}
df_pred = pd.DataFrame(data=q)
df_pred = df_pred[['ID', 'no_financial_services', 'other_only', 'mm_only', 'mm_plus']]

In [ ]: rfc.score(X_train, y_train)

```

0.6 XG BOOST

```

In [ ]: #Applying XGBoost
import xgboost as xgb

xgb_clf = xgb.XGBClassifier()
xgb_clf.fit(train_X, train_y)

print('The accuracy of the XGBoost classifier on training data is {:.2f}'.format(xgb_clf.score(train_X, train_y)))
print('The accuracy of the XGBoost classifier on test data is {:.2f}'.format(xgb_clf.score(test_X, test_y)))

In [ ]: xgb = xgb.XGBClassifier()
model.fit(X, y)
holdout_predictions = model.predict(holdout[columns])

In [ ]: holdout_ids = holdout["id"]
submission_df = {"id": holdout_ids,
                 "CLASS": holdout_predictions}
submission = pd.DataFrame(submission_df)
submission.to_csv("submission2.csv", index=False)

```



```
In [ ]:
```